



Build a Text-To-Speech system

Google AI ML Winter Camp, SHA, Jan 14 -- 18, 2019

Description

In this project, you will build a speech synthesis system using state-of-the-art techniques.

The speech synthesis system has two parts:

- Go from characters to magnitude spectrogram (char2feats)
- Go from spectrogram to actual audio (vocoder = phase reconstruction & super-resolution), called WaveRNN.

A team can tackle each task, or you can divide your team into two. At the end, the teams can combine their results to produce a full end-to-end system. Both parts can be implemented independently.

If you finish early, take a look at the “Bonus” task, where you would start building a real-time demo server.

Content

Foreword:

Remember to use `py_utils.HasShape`, and use the google documentation style. Shapes must always be documented. Also, write tests which check the shapes. It's also useful to use `py_utils.CheckNumerics` every so often. Write unit tests for all modules. Add summaries for waveforms, spectrograms, with utterance IDs so that you can go back and listen to them.

Day 1: Data preparation

Download and run the lingvo toolkit. Try the WMT16 system (follow the [instructions](#)), just to get warmed up.

Download the LJ [dataset](#). Parameterize for:

- char2feats
- waveRNN

Create a directory called `lingvo/tasks/tts`. You will create a set of tools similar to the ASR [data preparation](#).

Remember to install `sox`. You will need to create a tool similar to the ASR [one](#).



For char2feats, you will create `tfrecords` of `tf.Example` which contain the following:

- `uttid`: utterance ID. This will be useful for debugging. We will show this in the tensorboard summaries, a string.
- `transcript`: the text sentences, a string.
- `frames`: the spectrograms, a float tensor.
- Optionally, you could pre-tokenize the transcript into a tensor of integer IDs.

Note: there is a “bug” in the current feature extractor. You cannot change the context/stride. Just extract the frames and use `tf.reshape(frames, [b, -1, 80])`. (See ASR code.) The char2feats model will tokenize the `transcript` into graphemes (characters).

The WaveRNN needs to map from spectrograms to audio samples:

- `uttid`: utterance ID. Always keep that for debugging.
- `frames`: as above, the spectrograms.
- `audio_samples`: 24kHz audio samples, 0-1 floats.
- For debugging, write the associated text as `transcript`.

Write the `input_generator` for each. Run it in isolation to verify that you get what you expect, just like we did for the [ASR](#) input generator. Make sure that the scale of the input samples is done correctly, otherwise the energies will be very small.

Bonus: write a utility function that scans that tf records to find an audio sample and plays out the audio, for a given `uttid`.

Char2feats:

- Create encoder (stacked LSTMs)
- Create decoder (attention + stacked LSTMs)
- Create loss

You need to read the paper carefully and make sure the model converges. You can start with the machine translation (MT) [encoder](#). A reference implementation is [here](#). A good system diagram is [here](#).

Bonus: Can you try the Transformer encoder? Does it lead to better quality?

WaveRNN:

- Convert to 12kHz, 8-bit samples for development. This will train faster, with, of course, a loss in quality
- Create network & dual softmax loss
- Bonus: Create sparse network



The encoding is relatively [straightforward](#). The difficulty here is that it's hard to train.

Bonus:

To make a fully functional system, you'll need:

- Create the Inference Graph
- Efficient CPU implementation for sparse networks
- Demo / serving

It will be hard to have a real-time demo, but you can start work towards that goal during the course, and finish later.

References

[1] Tacotron2

[2] WaveRNN